

**ML LAB ASSIGNMENT 3 DIABETES KNN**

Implement K-Nearest Neighbors algorithm on diabetes.csv dataset. Compute confusion matrix, accuracy, error rate, precision and recall on the given dataset. Dataset link : <https://www.kaggle.com/datasets/abdallamahgoub/diabete>

```
# importing important packages
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
```

```
# Reading the data
data = pd.read_csv('/content/diabetes.csv')
```

Double-click (or enter) to edit

```
data.head()
```

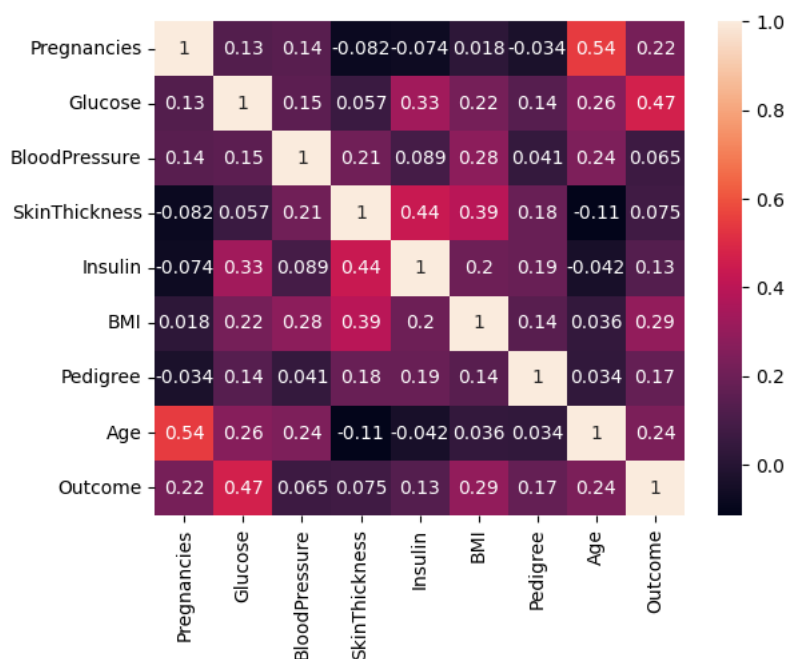
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	0
0	6	148	72	35	0	33.6	0.627	50	
1	1	85	66	29	0	26.6	0.351	31	
2	8	183	64	0	0	23.3	0.672	32	
3	1	89	66	23	94	28.1	0.167	21	
4	0	137	40	35	168	43.1	2.288	33	

```
# checking the Correlation between other columns and 'Outcome'
data.corr()['Outcome']
```

```
Pregnancies    0.221898
Glucose         0.466581
BloodPressure   0.065068
SkinThickness   0.074752
Insulin         0.130548
BMI             0.292695
Pedigree        0.173844
Age             0.238356
Outcome         1.000000
Name: Outcome, dtype: float64
```

```
# Heatmap shows relation between columns
sns.heatmap(data.corr(), annot=True)
```

<Axes: >



```

# Copying every column except 'Outcome'
X = data.drop('Outcome', axis=1)

# copying column "Outcome" to y df
y = data['Outcome']

# Importing train_test_split
from sklearn.model_selection import train_test_split

# Splitting the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

from sklearn.preprocessing import StandardScaler

from sklearn.neighbors import KNeighborsClassifier

# Creating object of StandardScaler()
scaler = StandardScaler()

# Creating object of KNN Classifier
knn = KNeighborsClassifier()

operations = [('scaler', scaler), ('knn', knn)]

from sklearn.pipeline import Pipeline

# Creating a pipeline instructions with given operations
# as initialised above
pipe = Pipeline(operations)

from sklearn.model_selection import GridSearchCV


k_values = list(range(1, 20))

# passing scoring metric
param_grid = {'knn__n_neighbors': k_values}

# passing hyperparameters, model and pipeline to GridSearchCV
# GridSearchCV will then train the model on all possible combinations of hyperparameter values
# and select the model that produces the best score on the scoring metric.
full_classifier = GridSearchCV(pipe, param_grid, cv=5, scoring='accuracy')

# Training the model
full_classifier.fit(X_train, y_train)

```



```

# Gives the hyperparameter values of the best model found by the grid search
full_classifier.best_estimator_.get_params()

```

```

{'memory': None,
 'steps': [('scaler', StandardScaler()),
            ('knn', KNeighborsClassifier(n_neighbors=7))],
 'verbose': False,
 'scaler': StandardScaler(),
 'knn': KNeighborsClassifier(n_neighbors=7),
 'scaler__copy': True,
 'scaler__with_mean': True,
 'scaler__with_std': True,
 'knn__algorithm': 'auto',
 'knn__leaf_size': 30,
 'knn__metric': 'minkowski',
 'knn__metric_params': None,
 'knn__n_jobs': None,
 'knn__n_neighbors': 7,
 'knn__p': 2,
 'knn__weights': 'uniform'}

```

```

# predicting the values
y_pred = full_classifier.predict(X_test)

from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

# Getting the accuracy score
accuracy_score(y_test, y_pred)

0.696969696969697

# Error rate is the percentage of predictions that are incorrect
# and calculated as follows :
# Error_rate = 1 - accuracy

print(1 - accuracy_score(y_test, y_pred))

0.303030303030303

# Getting confusion matrix
confusion_matrix(y_test, y_pred)

array([[121, 30],
       [ 40, 40]])

# Getting classification report
print(classification_report(y_test, y_pred))

              precision    recall  f1-score   support

     0           0.75         0.80         0.78         151
     1           0.57         0.50         0.53          80

 accuracy                   0.70         231
 macro avg           0.66         0.65         0.65         231
 weighted avg        0.69         0.70         0.69         231

from sklearn.metrics import accuracy_score, precision_score, recall_score

# Calculate the precision score
precision = precision_score(y_test, y_pred)

# Calculate the recall score
recall = recall_score(y_test, y_pred)

# Print the results
print("Precision:", precision)
print("Recall:", recall)

Precision: 0.5714285714285714
Recall: 0.5

```